

# Deep Learning Assignment 2: RNNs, Transformers, and Self-Supervised Learning

Due Date: **January 11th, 2026, 23:59**

## In this assignment:

- You will implement recurrent neural networks (RNNs) and Long Short-Term Memory (LSTM) networks from scratch.
- You will build an image captioning model using RNNs, LSTMs, and attention mechanisms.
- You will implement the Transformer architecture from the "Attention is All You Need" paper.
- You will explore self-supervised learning with SimCLR contrastive learning.
- You will work with state-of-the-art pretrained models: CLIP and DINO.

**Setup:** Same as Assignment 1 - use the free tier of Google Colab and upload the assignment folder to `My Drive/dl/assignments/assignment2/`. All notebooks require GPU acceleration for training.

## Assignment Parts

This assignment consists of four main parts. Complete all parts in the order presented.

### Part 1: Image Captioning with RNNs and LSTMs (25 points)

**Notebook:** `rnn_lstm_captioning.ipynb`

In this part, you will:

- Implement vanilla Recurrent Neural Networks (RNNs) from scratch
- Implement Long Short-Term Memory (LSTM) networks
- Build an image captioning model that generates natural language descriptions of images
- Implement attention mechanisms to improve caption quality
- Train your models on the COCO Captions dataset

You will implement RNN and LSTM cells with forward and backward passes, then combine them with a pre-trained CNN (RegNet-X 400MF) to create an end-to-end image captioning system. Finally, you'll augment your model with spatial attention to focus on relevant image regions while generating captions.

## Part 2: Transformers for Arithmetic Operations (25 points)

### Notebook: `Transformers.ipynb`

In this part, you will:

- Implement the Transformer architecture from "Attention is All You Need"
- Build core components: Multi-Head Attention, Layer Normalization, Feed-Forward blocks
- Implement both Encoder and Decoder blocks with residual connections
- Implement positional encoding (simple and sinusoidal)
- Train a Transformer model on a toy arithmetic dataset (addition and subtraction)
- Visualize attention weights to understand what the model learns

This implementation focuses on a simplified sequence-to-sequence task with fixed-length inputs and outputs, making it easier to understand the core Transformer mechanisms before applying them to more complex problems.

## Part 3: Self-Supervised Learning with SimCLR (25 points)

### Notebook: `Self_Supervised_Learning.ipynb`

In this part, you will:

- Learn about self-supervised learning and contrastive learning
- Implement data augmentation for SimCLR
- Implement the contrastive loss function (both naive and vectorized versions)
- Train a SimCLR model and compare it with a baseline model
- Observe the power of self-supervised pretraining for downstream classification tasks

SimCLR learns visual representations by maximizing agreement between differently augmented views of the same image. You will see how models pretrained with self-supervised learning outperform randomly initialized models on classification tasks.

## Part 4: CLIP and DINO (25 points)

### Notebook: `CLIP_DINO.ipynb`

In this part, you will:

- Explore CLIP (Contrastive Language-Image Pre-Training) for zero-shot classification
- Implement similarity computation between text and image features
- Build a zero-shot classifier and an image retrieval system using CLIP
- Explore DINO (self-DIstillation with NO labels) features for semantic segmentation
- Train a simple segmentation model on DINO features using the DAVIS dataset

CLIP learns to align image and text representations in a shared embedding space, enabling powerful zero-shot capabilities. DINO learns fine-grained visual features that are useful for dense prediction tasks like segmentation.

## Submission Instructions

When you have completed all parts of the assignment, follow these steps to generate and submit your work:

### Step 1: Generate Submission Files

1. Ensure all code cells in all notebooks have been executed and their outputs are visible
2. **IMPORTANT:** Manually save all `*.ipynb` and `*.py` files before proceeding
3. In the `CLIP_DINO.ipynb` notebook, navigate to the last cell
4. Run the submission cell, which will automatically generate:
  - `a2_code_submission.zip` – A ZIP archive containing:
    - `transformers.py`
    - `Transformers.ipynb`
    - `rnn_lstm_captioning.py`
    - `rnn_lstm_captioning.ipynb`
    - `Self_Supervised_Learning.ipynb`
    - `CLIP_DINO.ipynb`
    - `rnn_lstm_attention_submission.pt`
  - `a2_inline_submission.pdf` – A PDF containing all notebook outputs

### Step 2: Create Student Information File

Manually create a text file named `students.txt` containing both students' information in the following format:

`Student1_Full_Name Student1_ID`  
`Student2_Full_Name Student2_ID`

For example:

`John_Doe 123456789`  
`Jane_Smith 987654321`

**Note:** Use underscores (`_`) instead of spaces in the names, and ensure the student IDs match your official registration records.

### Step 3: Submit to Lemida

Download all files from your Google Drive to your local computer, then submit the following to Lemida:

1. `a2_inline_submission.pdf`
2. `a2_code_submission.zip`
3. `students.txt`

Good luck!